

Attention Is All You Need

Vaswani et al (2017)

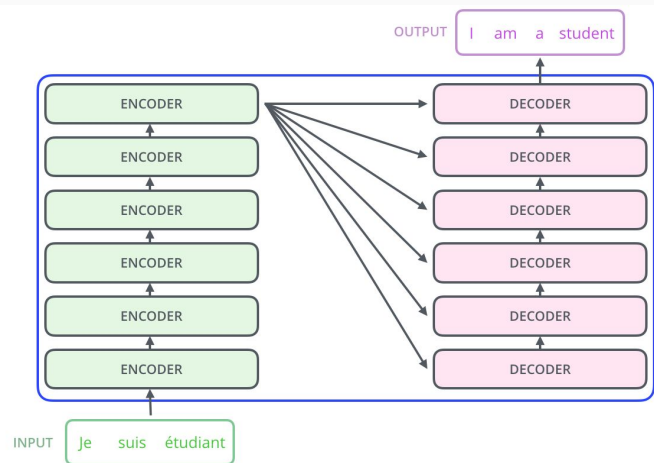
NLP reading club 2021/06/24

Key Points

The paper introduces the Transformer architecture, which:

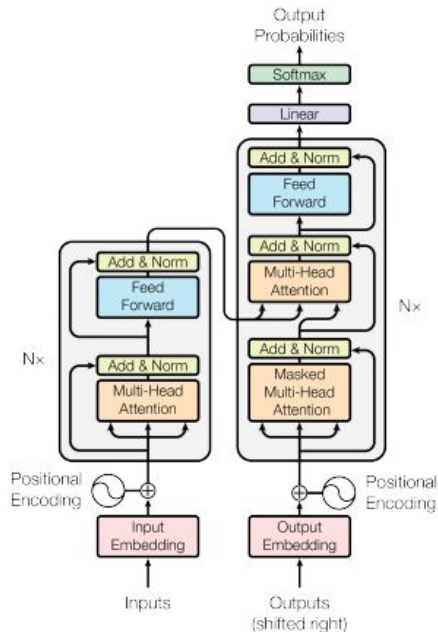
- Rely completely on attention to learn the sequence of input instead of using recurrent/convolutional units
- Uses multiheaded attention over the same input to get different projections to improve model performance
- Positional encodings to encode word ordering
- Faster training times as compared to RNNs/CNNs trained to do the same (translation/seq2seq) task

Model architecture



from <http://jalammar.github.io/illustrated-transformer/>

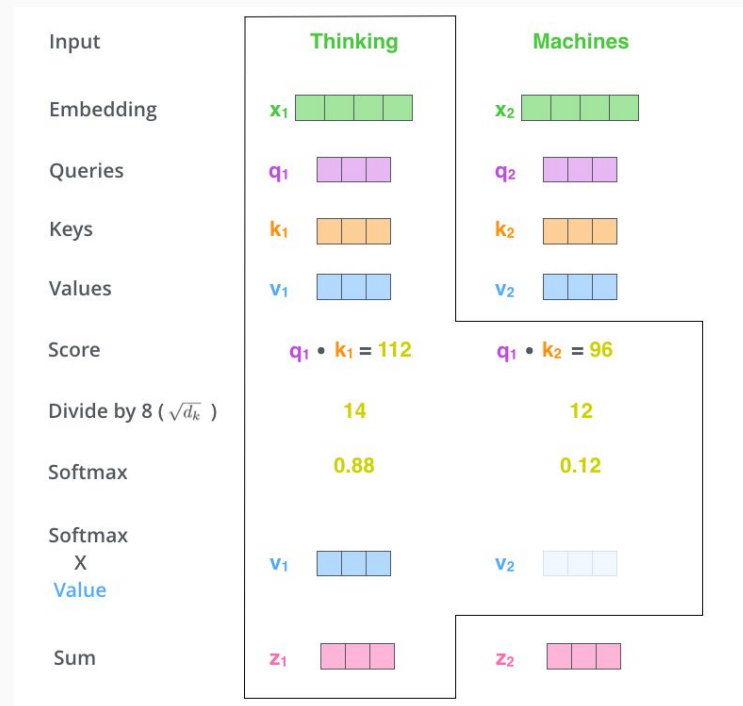
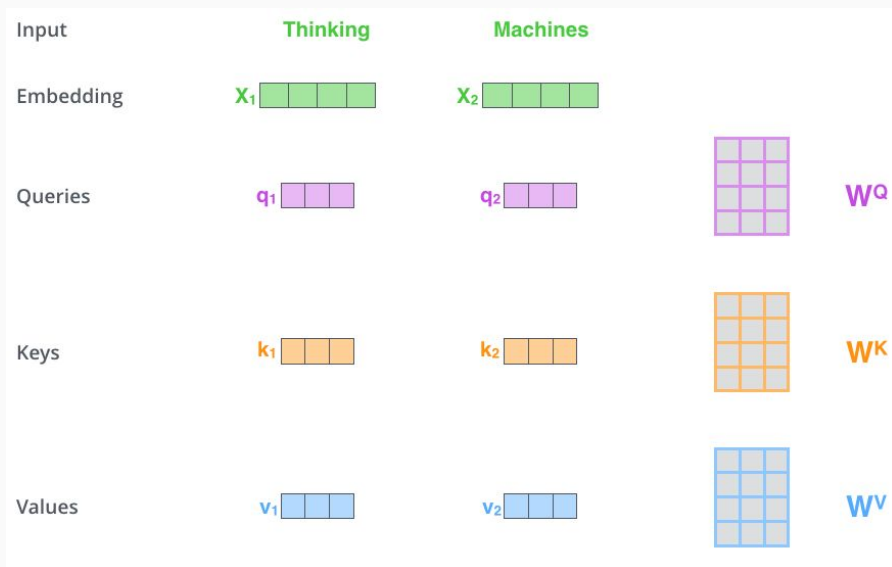
Similar to previous seq2seq models, the transformer has an encoder-decoder setup. However, multiple words can now pass through the model at once as opposed to RNN architectures where words have to be fed sequentially



Why self-attention

- Consider constant number of positions regardless of length of input (as opposed to RNNs, which have to consider $O(n)$ positions)
- Easily parallelizable
- Path length from distance part of the sentence is the same as those that are closer (unless you restrict the attention to a certain distance)

Attention mechanism



Multihead attention

1) This is our input sentence*

Thinking
Machines



2) We embed each word*




* In all encoders other than #0, we don't need embedding. We start directly with the output of the encoder right below this one




3) Split into 8 heads. We multiply X or R with weight matrices



...




4) Calculate attention using the resulting $Q/K/V$ matrices



...



5) Concatenate the resulting Z matrices, then multiply with weight matrix W^O to produce the output of the layer




...



W^O



Z



Experiment results

Table 2: The Transformer achieves better BLEU scores than previous state-of-the-art models on the English-to-German and English-to-French newstest2014 tests at a fraction of the training cost.

Model	BLEU		Training Cost (FLOPs)	
	EN-DE	EN-FR	EN-DE	EN-FR
ByteNet [18]	23.75			
Deep-Att + PosUnk [39]		39.2		$1.0 \cdot 10^{20}$
GNMT + RL [38]	24.6	39.92	$2.3 \cdot 10^{19}$	$1.4 \cdot 10^{20}$
ConvS2S [9]	25.16	40.46	$9.6 \cdot 10^{18}$	$1.5 \cdot 10^{20}$
MoE [32]	26.03	40.56	$2.0 \cdot 10^{19}$	$1.2 \cdot 10^{20}$
Deep-Att + PosUnk Ensemble [39]		40.4		$8.0 \cdot 10^{20}$
GNMT + RL Ensemble [38]	26.30	41.16	$1.8 \cdot 10^{20}$	$1.1 \cdot 10^{21}$
ConvS2S Ensemble [9]	26.36	41.29	$7.7 \cdot 10^{19}$	$1.2 \cdot 10^{21}$
Transformer (base model)	27.3	38.1	$3.3 \cdot 10^{18}$	
Transformer (big)	28.4	41.8	$2.3 \cdot 10^{19}$	

Experiment results

Parameter variations

	N	d_{model}	d_{ff}	h	d_k	d_v	P_{drop}	ϵ_{ls}	train steps	PPL (dev)	BLEU (dev)	params $\times 10^6$		
base	6	512	2048	8	64	64	0.1	0.1	100K	4.92	25.8	65		
(A)					1	512	512				5.29	24.9		
					4	128	128				5.00	25.5		
					16	32	32				4.91	25.8		
					32	16	16				5.01	25.4		
(B)					16					5.16	25.1	58		
					32					5.01	25.4	60		
(C)	2									6.11	23.7	36		
	4									5.19	25.3	50		
	8									4.88	25.5	80		
		256					32	32				5.75	24.5	28
		1024					128	128				4.66	26.0	168
			1024									5.12	25.4	53
			4096									4.75	26.2	90
(D)									0.0	5.77	24.6			
									0.2	4.95	25.5			
									0.0	4.67	25.3			
									0.2	5.47	25.7			
(E)	positional embedding instead of sinusoids									4.92	25.7			
big	6	1024	4096	16					0.3	300K	4.33	26.4	213	

Further resources

- [The illustrated transformer](#)
- [The Annotated Transformer](#)
- <https://github.com/xmu-xiaoma666/External-Attention-pytorch>